

Ridge Regression

Gregoire Gasparini, Aurora Hofman, Sarah Musiol, Beatriu Tort

25 de febrero de 2020

For the Assignment about Ridge Regression, we compute a function to choose a penalization parameter. The theory about Ridge Regression is used to write basic functions. In order to test these functions the Boston Housing data is used with it.

Alternatively, we proposed to use a package that deals with Ridge Regression in the background.

Choosing the penalization parameter

Function for choosing the penalization parameter

```
prostate <- read.table("prostate_data.txt", header=TRUE, row.names = 1)

train.sample <- which(prostate$train==TRUE) ##separate trainingsdata from testdata
val.sample <- which(prostate$train==FALSE)

Y_t <- scale( prostate$lpsa[train.sample], center=TRUE, scale=FALSE) ## center but not scale for respon
X_t <- scale( as.matrix(prostate[train.sample,1:8]), center=TRUE, scale=TRUE) ##scale and center for re
Y_val <- scale( prostate$lpsa[val.sample], center=TRUE, scale=FALSE) ## center but not scale for respon
X_val <- scale( as.matrix(prostate[val.sample,1:8]), center=TRUE, scale=TRUE) ##scale and center for re

#predictors

p <- dim(X_t)[2]
n <- dim(X_t)[1]

XtX <- t(X_t)%*%X_t

d2 <- eigen(XtX,symmetric = TRUE, only.values = TRUE)$values #eigenvalues of xtX

(cond.number <- sqrt(max(d2)/min(d2)))

## [1] 4.435608
lambda.max = 2e4

n_lambdas <- 25 ## look at 25 different values
```

```

lambda.v <- exp(seq(0,log(lambda.max+1),length=n_lambdas))-1 #lambda vector

n_val <- length(Y_val)

PMSE_vs <- function(X_t, Y_t, X_val, Y_val, lambda){

  p <- dim(X_t)[2]
  n_lambdas <- length(lambda)
  XtX <- t(X_t)%*%X_t
  PMSE_vec <- vector("numeric", length = n_lambdas)

  for(l in 1:n_lambdas){
    lambda <- lambda.v[l]
    beta_hat <- solve(XtX + lambda*diag(1,p)) %*% t(X_t) %*% Y_t
    #y_hat = X %*% beta_hat
    m_hat_vec <- vector("numeric", length = n_val)

    for (n in 1:n_val){
      m_hat_vec[n] <- (Y_val[n]-(X_val[n,]%*%beta_hat))^2
    }

    PMSE_vec[l]<- sum(m_hat_vec)/n_val

  }

  return (PMSE_vec)
}

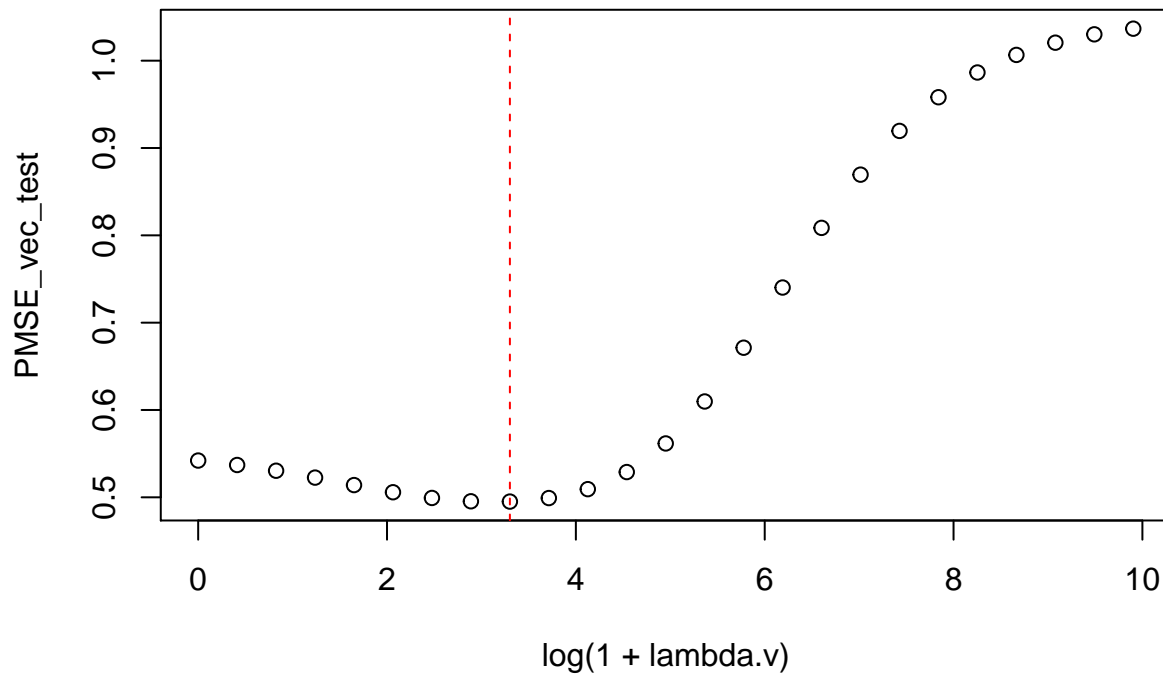
PMSE_vec_test <- PMSE_vs(X_t, Y_t, X_val, Y_val, lambda.v)

lambda.CV <- lambda.v[which.min(PMSE_vec_test)]

plot(log(1+lambda.v), PMSE_vec_test)

abline(v=log(1+lambda.CV),col=2,lty=2)

```



Function for Implementing the Ridge Regression penalization parameter

```
PMSE_k_fold <- function(X_t, Y_t, lambda.v, k=10){

  n_X_t <- dim(X_t)[1]
  p_X_t <- dim(X_t)[2] #length of columns
  n_subset <- as.integer((n_X_t)/k)+ 1
  group <- rep(seq(1,k), times = n_subset)
  group <- group[1:n_X_t]
group_random <- sample(group)
X_t_group <- cbind(X_t, group_random)
Y_t_group <- cbind(Y_t, group_random)

#now we can start:

PMSE <- list()

for (la in 1:n_lambdas){

  group_random <- sample(group)

  X_t_group <- cbind(X_t, group_random)

  Y_t_group <- cbind(Y_t, group_random)

  #now we can start:

  PMSE <- list()

  for (la in 1:n_lambdas){
```

```

lambda <- lambda.v[la]

y_hat <- list()
beta <- list()
h <- list()
y <- list()

for (l in 1:k){

  new_X_t_val <- subset(X_t_group, group_random==1)[ ,1:p_X_t]
  new_X_t_test <- subset(X_t_group, group_random!=1)[ ,1:p_X_t]

  new_Y_t_val <- subset.matrix(Y_t_group, group_random==1)[,1]
  new_Y_t_test <- subset.matrix(Y_t_group, group_random!=1)[,1]

  p_new <- dim(new_X_t_test)[2]
  p_new_v <- dim(new_X_t_val)[2]

  beta[[l]] <- solve(t(new_X_t_test)%*%new_X_t_test + lambda*diag(1,p_new))%*% t(new_X_t_test)%*%(n
  H_val <- new_X_t_val%*%solve(t(new_X_t_val)%*%new_X_t_val + (lambda+1e-13)*diag(1,p_new_v))%*% t(n

  # singular matrix for lambda = 0 -> trick: add a very small number

  y_hat[[l]] <- (new_X_t_val)%*%beta[[l]]
  h[[l]] <- diag(H_val)
  y[[l]] <- new_Y_t_val

}

y_hat <- c(do.call(rbind, y_hat))
beta <- c(do.call(cbind, beta))
h <- c(do.call(rbind, h))
y <- c(do.call(rbind, y))

PMSE[[la]] <- 1/n_X_t * sum(((y-y_hat)/(1-h))^2)

}

}

return(PMSE)


}

PMSE <- PMSE_k_fold(X_t = X_t, Y_t = Y_t, lambda.v = lambda.v, k = 10)

plot(log(lambda.v[-1]+1), PMSE[-1], ylab = "Predicted Mean Squarred Error", xlab = "log(Lambda+1)", mai

```

10-fold of Training set




In the graphic of the PMSE of 10-fold of Training set we have a negative exponential distribution. We suspect that there is some problem in the code and this is not working normally but we weren't able to find it. Theoretically we should have a graphic with a lower scale and that goes down and then up again so we have a minimum lambda that would be the optimum one.

Comparison between 5-fold and 10-fold

```
PMSE_val_5 <- PMSE_k_fold(X_t = X_val, Y_t = Y_val, lambda.v = lambda.v, k = 5)


plot(log(lambda.v[-1]+1), PMSE_val_5[-1], ylab = "Predicted Mean Squarred Error", xlab = "log(Lambda+1)")
```

5-fold of Validation set




```
PMSE_val_10 <- PMSE_k_fold(X_t = X_val, Y_t = Y_val, lambda.v = lambda.v, k = 10)
plot(log(lambda.v[-1]+1), PMSE_val_10[-1], ylab = "Predicted Mean Squarred Error", xlab = "log(Lambda+1)
```

10-fold of Validation set



#I am not sure about if the one below is the leave-one-out cross validation or not.

```
plot(log(1+lambda.v), PMSE_vec_test)
abline(v=log(1+lambda.CV), col=2, lty=2)
```

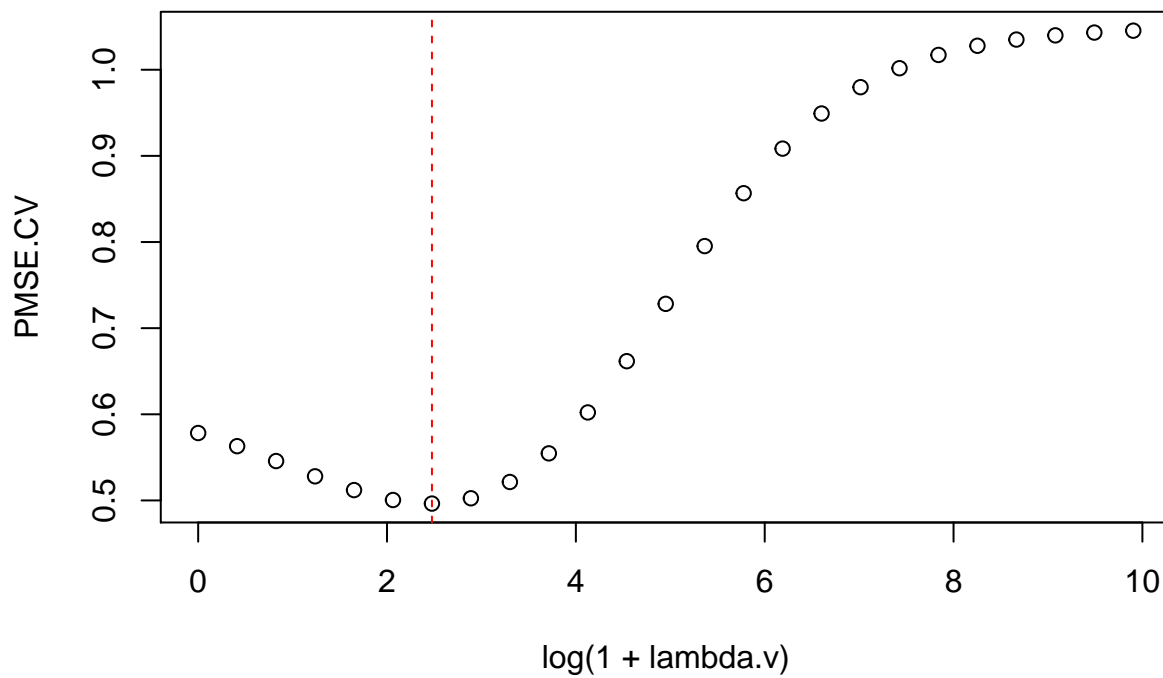


```
## leave-one-out cross validation

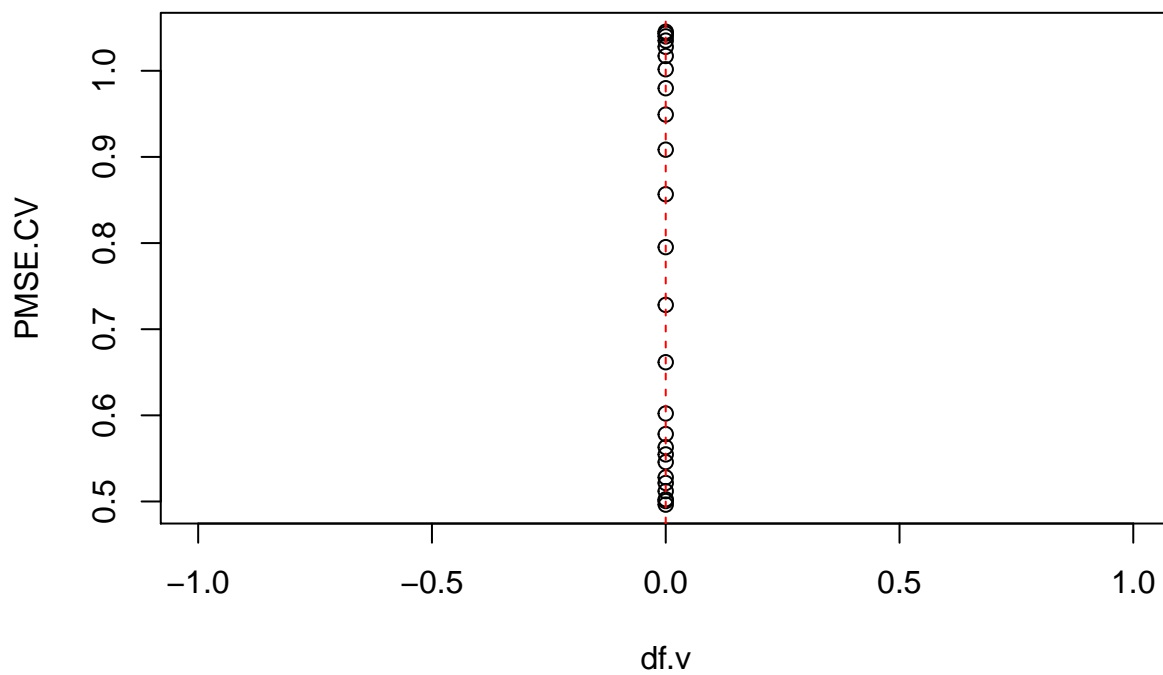
df.v <- numeric(n_lambdas)

PMSE.CV <- numeric(n_lambdas)
for (l in 1:n_lambdas){
  lambda <- lambda.v[l]
  PMSE.CV[l] <- 0
  for (i in 1:n_val){
    # m.Y.i <- mean(Y[-i])
    m.Y.i <- 0
    X.i <- X_val[-i,]
    Y.i <- Y_val[-i] - m.Y.i
    Xi <- X_val[i,]
    Yi <- Y_val[i]
    beta.i <- solve(t(X.i)%*%X.i + lambda*diag(1,p)) %*% t(X.i) %*% Y.i
    hat.Yi <- Xi %*% beta.i + m.Y.i
    PMSE.CV[l] <- PMSE.CV[l] + (hat.Yi - Yi)^2
  }
  PMSE.CV[l] <- PMSE.CV[l]/n_val
}
lambda.CV <- lambda.v[which.min(PMSE.CV)]
df.CV <- df.v[which.min(PMSE.CV)]

plot(log(1+lambda.v), PMSE.CV)
abline(v=log(1+lambda.CV), col=2, lty=2)
```



```
plot(df.v, PMSE.CV)
abline(v=df.CV,col=2,lty=2)
```



```
## Generalized Cross Validation (GCV)

beta.path <- matrix(0,nrow=n_lambdas, ncol=p)
diag.H.lambda <- matrix(0,nrow=n_lambdas, ncol=n_val)

PMSE.GCV <- numeric(n_lambdas)
for (l in 1:n_lambdas){
```



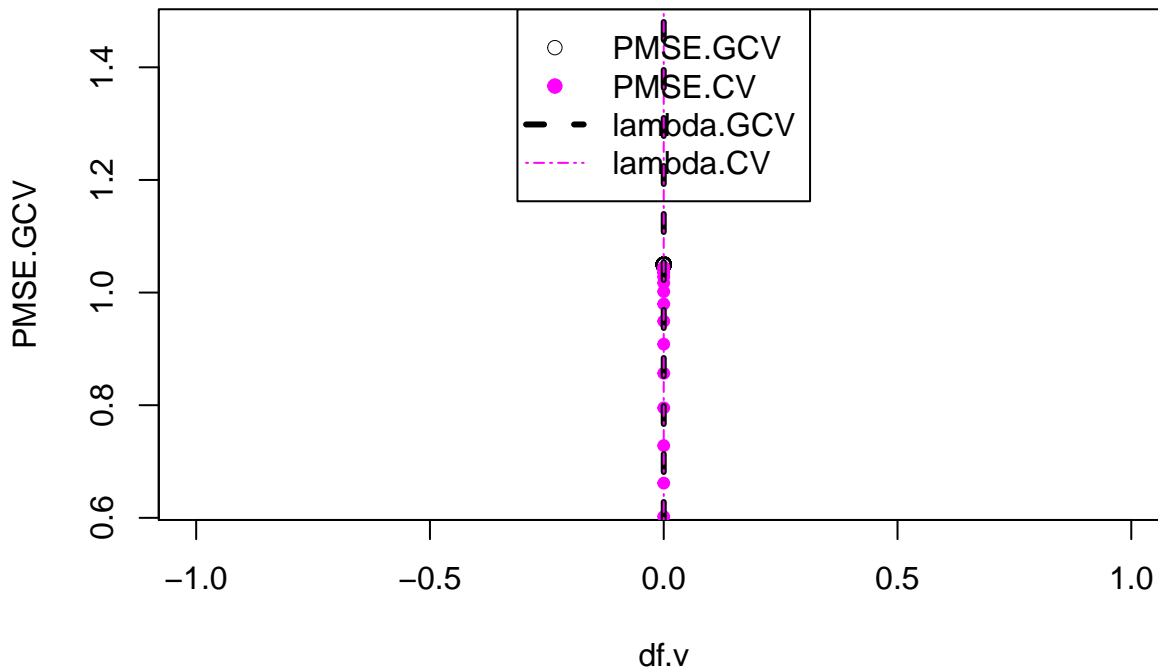
```

lambda <- lambda.v[1]
hat.Y <- X_val %*% beta.path[1,]
nu <- sum(diag.H.lambda[1,])
PMSE.GCV[1] <- sum( ((Y_val-hat.Y)/(1-nu/n_val))^2 )/n_val
}
lambda.GCV <- lambda.v[which.min(PMSE.GCV)]
df.GCV <- df.v[which.min(PMSE.GCV)]

PMSE.CV.H.lambda <- numeric(n_lambdas)
lambda.CV.H.lambda <- lambda.v[which.min(PMSE.CV.H.lambda)]
df.CV.H.lambda <- df.v[which.min(PMSE.CV.H.lambda)]

plot(df.v, PMSE.GCV)
points(df.v, PMSE.CV,col=6,pch=19,cex=.75)
abline(v=df.GCV,col=1,lty=2,lwd=3)
abline(v=df.CV.H.lambda,col=6,lty=6)
legend("top",c("PMSE.GCV","PMSE.CV","lambda.GCV","lambda.CV"),
      pch=c(1,19,NA,NA),lty=c(0,0,2,6),lwd=c(0,0,3,1),col=c(1,6,1,6))

```



Ridge Regression for the Boston Housing data

Loading the (corrected) Boston Housing data

```

library(MASS)

data(Boston)

help(Boston)

```

```

boston <- load("boston.Rdata")

boston <- boston.c

# cv.glmnet cannot handel factors -> "CHAS" is a factor
boston.c$CHAS <- as.numeric(boston.c$CHAS)

response <- "MEDV"

explanatory <- c("CRIM", "ZN", "INDUS", "CHAS", "NOX", "RM", "AGE", "DIS", "RAD", "TAX", "PTRATIO", "B"

```

Using obtained functions for Boston Housing data


```

PMSE_10 <- PMSE_k_fold(X_t = as.matrix(boston.c[, explanatory]), Y_t = boston.c$MEDV, lambda.v = lambda.v, k = 10)

plot(log(lambda.v[-1]+1), PMSE_10[-1], ylab = "Predicted Mean Squarred Error", xlab = "log(Lambda+1)", r

```

10-fold for Boston Housing Data



Alternative solution for Ridge Regression for the Boston Housing data

There exists a package called 'glmnet' that deals with elastic nets. Specifying $\alpha = 0$ Ridge Regression is applied on the data.

```
#install.packages("glmnet")
```

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 3.0-2
```

```
(ridge <- glmnet(y = boston.c$MEDV, x = as.matrix(boston.c[, explanatory]), alpha = 0))
```

```
##
```

```
## Call:  glmnet(x = as.matrix(boston.c[, explanatory]), y = boston.c$MEDV,      alpha = 0)
```

```
##
```

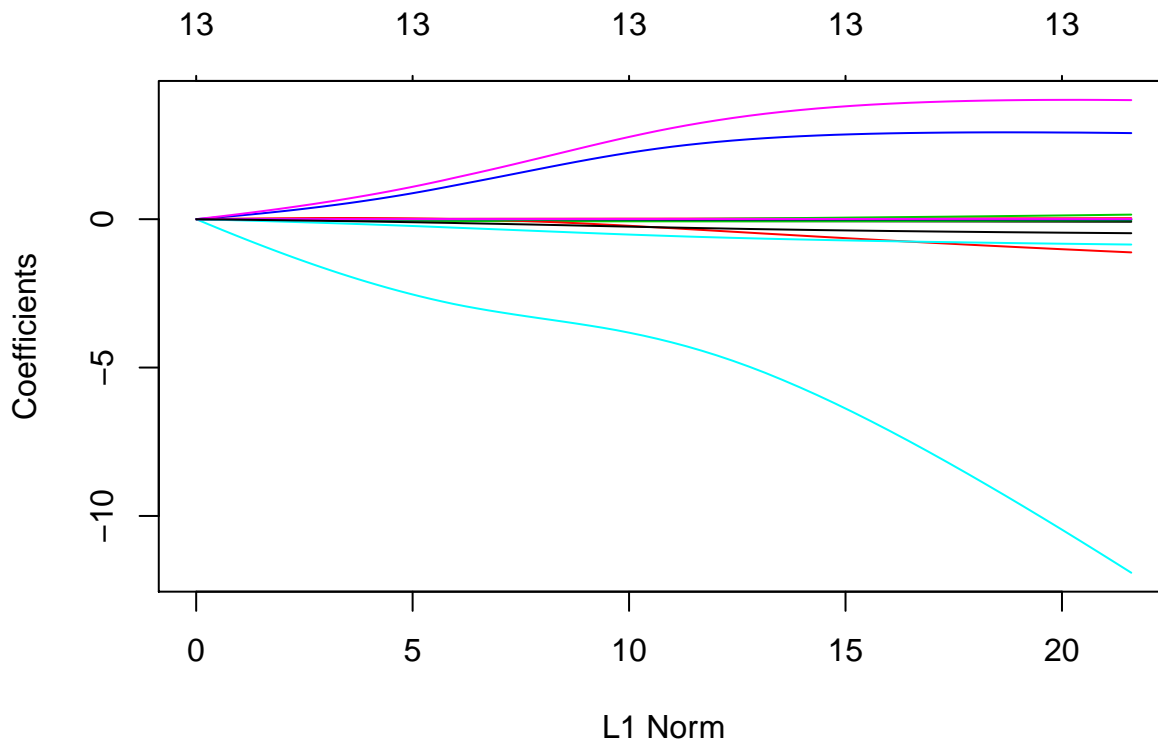
##		Df	%Dev	Lambda
## 1	13	0.00000	6778.0	
## 2	13	0.00792	6176.0	
## 3	13	0.00868	5627.0	
## 4	13	0.00952	5127.0	
## 5	13	0.01043	4672.0	
## 6	13	0.01143	4257.0	
## 7	13	0.01253	3878.0	
## 8	13	0.01372	3534.0	
## 9	13	0.01503	3220.0	
## 10	13	0.01646	2934.0	
## 11	13	0.01802	2673.0	
## 12	13	0.01972	2436.0	
## 13	13	0.02159	2219.0	
## 14	13	0.02362	2022.0	
## 15	13	0.02583	1843.0	
## 16	13	0.02824	1679.0	
## 17	13	0.03087	1530.0	
## 18	13	0.03372	1394.0	
## 19	13	0.03683	1270.0	
## 20	13	0.04020	1157.0	
## 21	13	0.04386	1054.0	
## 22	13	0.04783	960.7	
## 23	13	0.05213	875.4	
## 24	13	0.05678	797.6	
## 25	13	0.06180	726.7	
## 26	13	0.06721	662.2	
## 27	13	0.07304	603.4	
## 28	13	0.07932	549.8	
## 29	13	0.08605	500.9	
## 30	13	0.09326	456.4	
## 31	13	0.10100	415.9	
## 32	13	0.10920	378.9	
## 33	13	0.11790	345.3	
## 34	13	0.12720	314.6	
## 35	13	0.13710	286.6	
## 36	13	0.14750	261.2	
## 37	13	0.15840	238.0	
## 38	13	0.16990	216.8	
## 39	13	0.18200	197.6	
## 40	13	0.19450	180.0	
## 41	13	0.20760	164.0	
## 42	13	0.22110	149.5	
## 43	13	0.23510	136.2	
## 44	13	0.24940	124.1	
## 45	13	0.26420	113.1	
## 46	13	0.27920	103.0	
## 47	13	0.29450	93.9	

## 48	13	0.31000	85.5
## 49	13	0.32570	77.9
## 50	13	0.34150	71.0
## 51	13	0.35730	64.7
## 52	13	0.37320	59.0
## 53	13	0.38900	53.7
## 54	13	0.40470	48.9
## 55	13	0.42030	44.6
## 56	13	0.43570	40.6
## 57	13	0.45090	37.0
## 58	13	0.46590	33.7
## 59	13	0.48060	30.7
## 60	13	0.49490	28.0
## 61	13	0.50900	25.5
## 62	13	0.52270	23.2
## 63	13	0.53600	21.2
## 64	13	0.54890	19.3
## 65	13	0.56140	17.6
## 66	13	0.57340	16.0
## 67	13	0.58500	14.6
## 68	13	0.59610	13.3
## 69	13	0.60660	12.1
## 70	13	0.61670	11.1
## 71	13	0.62620	10.1
## 72	13	0.63530	9.2
## 73	13	0.64380	8.4
## 74	13	0.65170	7.6
## 75	13	0.65920	6.9
## 76	13	0.66610	6.3
## 77	13	0.67260	5.8
## 78	13	0.67860	5.2
## 79	13	0.68410	4.8
## 80	13	0.68920	4.4
## 81	13	0.69390	4.0
## 82	13	0.69820	3.6
## 83	13	0.70210	3.3
## 84	13	0.70570	3.0
## 85	13	0.70900	2.7
## 86	13	0.71200	2.5
## 87	13	0.71480	2.3
## 88	13	0.71730	2.1
## 89	13	0.71960	1.9
## 90	13	0.72170	1.7
## 91	13	0.72360	1.6
## 92	13	0.72530	1.4
## 93	13	0.72690	1.3
## 94	13	0.72830	1.2
## 95	13	0.72960	1.1
## 96	13	0.73080	1.0
## 97	13	0.73190	0.9
## 98	13	0.73280	0.8
## 99	13	0.73370	0.7
## 100	13	0.73450	0.7

```
# alpha = 0: Ridge Regression
```

```
# alpha = 1: Lasso Regression
```

```
plot(ridge)
```



```
(cv.ridge <- cv.glmnet(y = boston.c$MEDV, x = as.matrix(boston.c[, explanatory]), alpha = 0))
```

```
##
## Call:  cv.glmnet(x = as.matrix(boston.c[, explanatory]), y = boston.c$MEDV,      alpha = 0)
##
## Measure: Mean-Squared Error
##
##      Lambda Measure      SE Nonzero
## min  0.678    24.09 3.567      13
## 1se  4.357    27.50 4.025      13
```

```
plot(cv.ridge)
```

