

Ridge Regression

Gregoire Gasparini, Aurora Hofman, Sarah Musiol, Beatriu Tort

25/02/20

For the Assignment about Ridge Regression, we compute a function to choose a penalization parameter. The theory about Ridge Regression is used to write basic functions. In order to test these functions the Boston Housing data is used with it. Alternatively, we proposed to use a package that deals with Ridge Regression in the background.

Choosing the penalization parameter

Function for choosing the penalization parameter

```
prostate <- read.table("prostate_data.txt", header=TRUE, row.names = 1)
train.sample <- which(prostate$train==TRUE) ##separate trainingsdata from testdata
val.sample <- which(prostate$train==FALSE)

Y_t <- scale(prostate$lpsa[train.sample], center=TRUE, scale=FALSE) ## center but not scale for respon.
X_t <- scale(as.matrix(prostate[train.sample,1:8]), center=TRUE, scale=TRUE) ##scale and center for

Y_val <- scale(prostate$lpsa[val.sample], center=TRUE, scale=FALSE) ## center but not scale for respon.
X_val <- scale(as.matrix(prostate[val.sample,1:8]), center=TRUE, scale=TRUE)

#predictors
p <- dim(X_t)[2]

XtX <- t(X_t)%*%X_t
d2 <- eigen(XtX,symmetric = TRUE, only.values = TRUE)$values #eigenvalues of xtx

(cond.number <- sqrt(max(d2)/min(d2)))

## [1] 4.435608
lambda.max = 2e4
n_lambdas <- 25 ## look at 25 different values
lambda.v <- exp(seq(0,log(lambda.max+1),length=n_lambdas))-1 #lambda vector

n_val <- length(Y_val)

PMSE_vs <- function(X_t, Y_t, X_val, Y_val, lambda){
  p <- dim(X_t)[2]
  n_lambdas <- length(lambda)

  XtX <- t(X_t)%*%X_t
```

```

PMSE_vec <- vector("numeric", length = n_lambdas)
for(l in 1:n_lambdas){
  lambda <- lambda.v[l]
  beta_hat <- solve(XtX + lambda*diag(1,p)) %*% t(X_t) %*% Y_t
  #y_hat = X %*% beta_hat
  m_hat_vec <- vector("numeric", length = n_val)

  for (n in 1:n_val){
    m_hat_vec[n] <- (Y_val[n]-(X_val[n,]%*%beta_hat))^2
  }
  PMSE_vec[l]<- sum(m_hat_vec)/n_val
}
return (PMSE_vec)
}


```

```
PMSE_vec_test <- PMSE_vs(X_t, Y_t, X_val, Y_val, lambda.v)
```

```

lambda.CV <- lambda.v[which.min(PMSE_vec_test)]
plot(log(1+lambda.v), PMSE_vec_test)
abline(v=log(1+lambda.CV),col=2,lty=2)

```



Function for Implementing the Ridge Regression penalization parameter

```

PMSE_k_fold <- function(X_t, Y_t, lambda.v, k=10){
  n_X_t <- dim(X_t)[1]
  p_X_t <- dim(X_t)[2] #length of columns
  n_subset <- as.integer((n_X_t)/k)+ 1
  group <- rep(seq(1,k), times = n_subset)
  group <- group[1:n_X_t]
}

```

```

group_random <- sample(group)

X_t_group <- cbind(X_t, group_random)
Y_t_group <- cbind(Y_t, group_random)

#now we can start:
PMSE <- list()

for (la in 1:n_lambdas){

  group_random <- sample(group)

  X_t_group <- cbind(X_t, group_random)
  Y_t_group <- cbind(Y_t, group_random)

  #now we can start:
  PMSE <- list()

  for (la in 1:n_lambdas){

    lambda <- lambda.v[la]

    y_hat <- list()
    beta <- list()
    h <- list()
    y <- list()

    for (l in 1:k){
      new_X_t_val <- subset(X_t_group, group_random==1)[,1:p_X_t]
      new_X_t_test <- subset(X_t_group, group_random!=1)[,1:p_X_t]

      new_Y_t_val <- subset.matrix(Y_t_group, group_random==1)[,1]
      new_Y_t_test <- subset.matrix(Y_t_group, group_random!=1)[,1]

      p_new <- dim(new_X_t_test)[2]
      p_new_v <- dim(new_X_t_val)[2]

      beta[[l]] <- solve(t(new_X_t_test)%*%new_X_t_test + lambda*diag(1,p_new))%*% t(new_X_t_test)%*%(n

      H_val <- new_X_t_val%*%solve(t(new_X_t_val)%*%new_X_t_val + (lambda+1e-13)*diag(1,p_new_v))%*% t(
      # singular matrix for lambda = 0 -> trick: add a very small number
      y_hat[[l]] <- (new_X_t_val)%*%beta[[l]]
      h[[l]] <- diag(H_val)
      y[[l]] <- new_Y_t_val

    }

    y_hat <- c(do.call(rbind, y_hat))
    beta <- c(do.call(cbind, beta))
    h <- c(do.call(rbind, h))
  }
}

```


```

y <- c(do.call(rbind, y))

PMSE[[1a]] <- 1/n_X_t * sum(((y-y_hat)/(1-h))^2)
}
}
return(PMSE)
}

PMSE <- PMSE_k_fold(X_t = X_t, Y_t = Y_t, lambda.v = lambda.v, k = 10)
plot(log(lambda.v[-1]+1), PMSE[-1], ylab = "Predicted Mean Squarred Error", xlab = "log(Lambda+1)")

```




Comparison between 5-fold and 10-fold


```

PMSE_val_5 <- PMSE_k_fold(X_t = X_val, Y_t = Y_val, lambda.v = lambda.v, k = 5)
plot(log(lambda.v[-1]+1), PMSE_val_5[-1], ylab = "Predicted Mean Squarred Error", xlab = "log(Lambda+1)")

```



```
PMSE_val_10 <- PMSE_k_fold(X_t = X_val, Y_t = Y_val, lambda.v = lambda.v, k = 10)
plot(log(lambda.v[-1]+1), PMSE_val_10[-1], ylab = "Predicted Mean Squarred Error", xlab = "log(Lambda+1)")
```



Ridge Regression for the Boston Housing data

Loading the (corrected) Boston Housing data

```
library(MASS)
data(Boston)
help(Boston)
```

```
boston <- load("boston.Rdata")
```

Using obtained functions for Boston Housing data

Alternative solution for Ridge Regression for the Boston Housing data

There exists a package called 'glmnet' that deals with elastic nets. Specifying $\alpha = 0$ Ridge Regression is applied on the data.

```
#install.packages("glmnet")
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 3.0-2
```

```
response <- "MEDV"
```

```
explanatory <- c("CRIM", "ZN", "INDUS", "CHAS", "NOX", "RM", "AGE", "DIS", "RAD", "TAX", "PTRATIO", "B"
```

```
# cv.glmnet cannot handel factors -> "CHAS" is a factor
```

```
boston.c$CHAS <- as.numeric(boston.c$CHAS)
```

```
(ridge <- glmnet(y = boston.c$MEDV, x = as.matrix(boston.c[, explanatory]), alpha = 0))
```

```
##
```

```
## Call: glmnet(x = as.matrix(boston.c[, explanatory]), y = boston.c$MEDV, alpha = 0)
```

```
##
```

```
##      Df      %Dev Lambda
```

```
## 1    13 0.00000 6778.0
```

```
## 2    13 0.00792 6176.0
```

```
## 3    13 0.00868 5627.0
```

```
## 4    13 0.00952 5127.0
```

```
## 5    13 0.01043 4672.0
```

```
## 6    13 0.01143 4257.0
```

```
## 7    13 0.01253 3878.0
```

```
## 8    13 0.01372 3534.0
```

```
## 9    13 0.01503 3220.0
```

```
## 10   13 0.01646 2934.0
```

```
## 11   13 0.01802 2673.0
```

```
## 12   13 0.01972 2436.0
```

```
## 13   13 0.02159 2219.0
```

```
## 14   13 0.02362 2022.0
```

```
## 15   13 0.02583 1843.0
```

```
## 16   13 0.02824 1679.0
```

```
## 17   13 0.03087 1530.0
```

```
## 18   13 0.03372 1394.0
```

```
## 19   13 0.03683 1270.0
```

```
## 20   13 0.04020 1157.0
```

```
## 21   13 0.04386 1054.0
```

```
## 22   13 0.04783  960.7
```

```
## 23   13 0.05213  875.4
```

```
## 24   13 0.05678  797.6
```


```
## 25   13 0.06180  726.7
```

```
## 26   13 0.06721  662.2
```

##	27	13	0.07304	603.4
##	28	13	0.07932	549.8
##	29	13	0.08605	500.9
##	30	13	0.09326	456.4
##	31	13	0.10100	415.9
##	32	13	0.10920	378.9
##	33	13	0.11790	345.3
##	34	13	0.12720	314.6
##	35	13	0.13710	286.6
##	36	13	0.14750	261.2
##	37	13	0.15840	238.0
##	38	13	0.16990	216.8
##	39	13	0.18200	197.6
##	40	13	0.19450	180.0
##	41	13	0.20760	164.0
##	42	13	0.22110	149.5
##	43	13	0.23510	136.2
##	44	13	0.24940	124.1
##	45	13	0.26420	113.1
##	46	13	0.27920	103.0
##	47	13	0.29450	93.9
##	48	13	0.31000	85.5
##	49	13	0.32570	77.9
##	50	13	0.34150	71.0
##	51	13	0.35730	64.7
##	52	13	0.37320	59.0
##	53	13	0.38900	53.7
##	54	13	0.40470	48.9
##	55	13	0.42030	44.6
##	56	13	0.43570	40.6
##	57	13	0.45090	37.0
##	58	13	0.46590	33.7
##	59	13	0.48060	30.7
##	60	13	0.49490	28.0
##	61	13	0.50900	25.5
##	62	13	0.52270	23.2
##	63	13	0.53600	21.2
##	64	13	0.54890	19.3
##	65	13	0.56140	17.6
##	66	13	0.57340	16.0
##	67	13	0.58500	14.6
##	68	13	0.59610	13.3
##	69	13	0.60660	12.1
##	70	13	0.61670	11.1
##	71	13	0.62620	10.1
##	72	13	0.63530	9.2
##	73	13	0.64380	8.4
##	74	13	0.65170	7.6
##	75	13	0.65920	6.9
##	76	13	0.66610	6.3
##	77	13	0.67260	5.8
##	78	13	0.67860	5.2
##	79	13	0.68410	4.8
##	80	13	0.68920	4.4

```
## 81 13 0.69390 4.0
## 82 13 0.69820 3.6
## 83 13 0.70210 3.3
## 84 13 0.70570 3.0
## 85 13 0.70900 2.7
## 86 13 0.71200 2.5
## 87 13 0.71480 2.3
## 88 13 0.71730 2.1
## 89 13 0.71960 1.9
## 90 13 0.72170 1.7
## 91 13 0.72360 1.6
## 92 13 0.72530 1.4
## 93 13 0.72690 1.3
## 94 13 0.72830 1.2
## 95 13 0.72960 1.1
## 96 13 0.73080 1.0
## 97 13 0.73190 0.9
## 98 13 0.73280 0.8
## 99 13 0.73370 0.7
## 100 13 0.73450 0.7
```

```
# alpha = 0: Ridge Regression
# alpha = 1: Lasso Regression
plot(ridge)
```



```
(cv.ridge <- cv.glmnet(y = boston.c$MEDV, x = as.matrix(boston.c[, explanatory]), alpha = 0))
```

```
##
## Call: cv.glmnet(x = as.matrix(boston.c[, explanatory]), y = boston.c$MEDV,      alpha = 0)
##
## Measure: Mean-Squared Error
```



```
##
##      Lambda Measure      SE Nonzero
## min  0.678    24.20  2.830      13
## 1se  3.617    26.83  3.397      13
```

```
plot(cv.ridge)
```

